

# Today in Cryptography (5830)

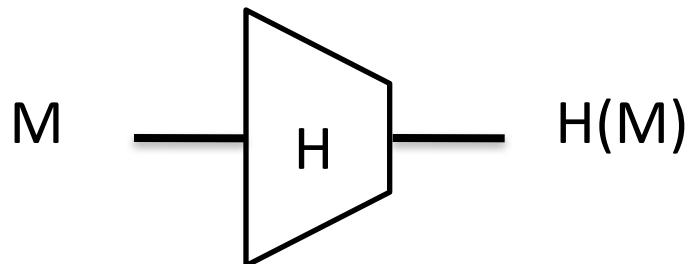
Password hashing  
Password-based AE

# Where we are at

- Authenticated encryption
  - Symmetric encryption providing confidentiality and integrity
- Hash functions
  - Collision resistance
  - Use as PRFs, MACs (HMAC)
- Today:
  - Password-based key derivation
  - Password-based AE

# Cryptographic hash functions

A function  $H$  that maps arbitrary bit string to fixed length string of size  $n$



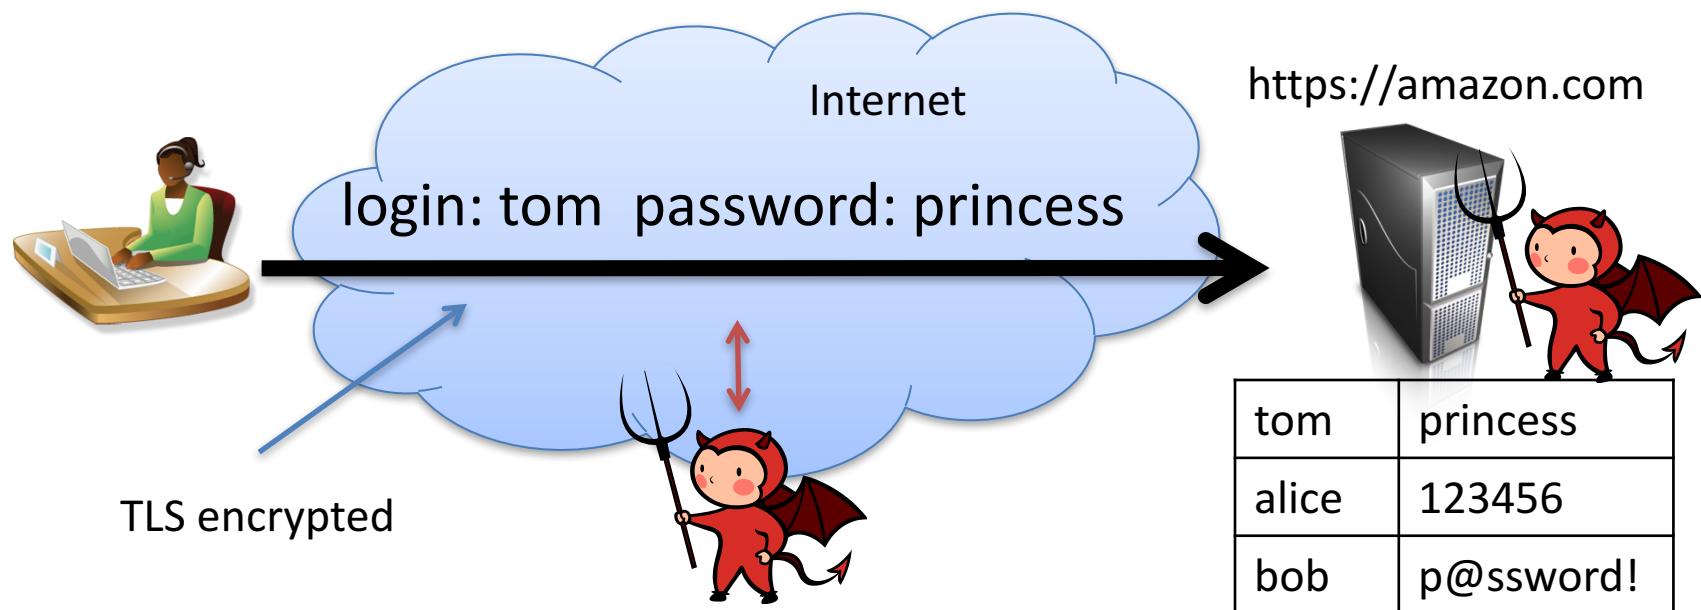
SHA-256:  $n = 256$  bits  
SHA-512:  $n = 512$  bits  
SHA-3:  $n = 224, 256, 384, 512$

Many security goals asked of hash functions. Ideally, they behave as if they were a (public) random function.

Security goals:

- Collision resistance
- Preimage resistance (one-wayness)
- Good as a PRF if we key it appropriately
- “Behave like” a public, random function

# Passwords



# Breaches are ubiquitous

:



32.6 million leaked (2012)  
32.6 million recovered (plaintext!)



6.5 million leaked (2012)  
5.85 million recovered in 2 weeks (SHA-1)



36 million accounts leaked (2013)  
Encrypted, but with ECB mode

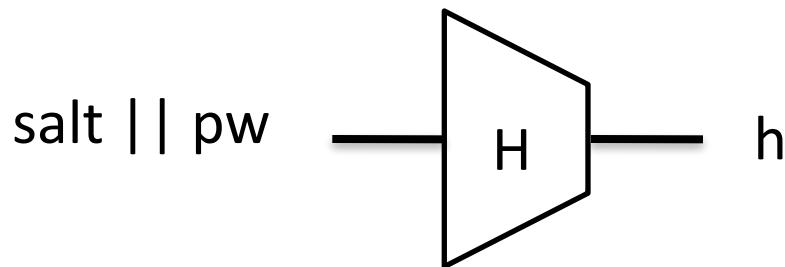


:

1 billion accounts (2013)  
MD5 hashes

# Password hashing

Password hashing. Choose random salt and store (salt,h) where:



**The idea:** Attacker, given (salt,h), should not be able to recover pw

Or can they?

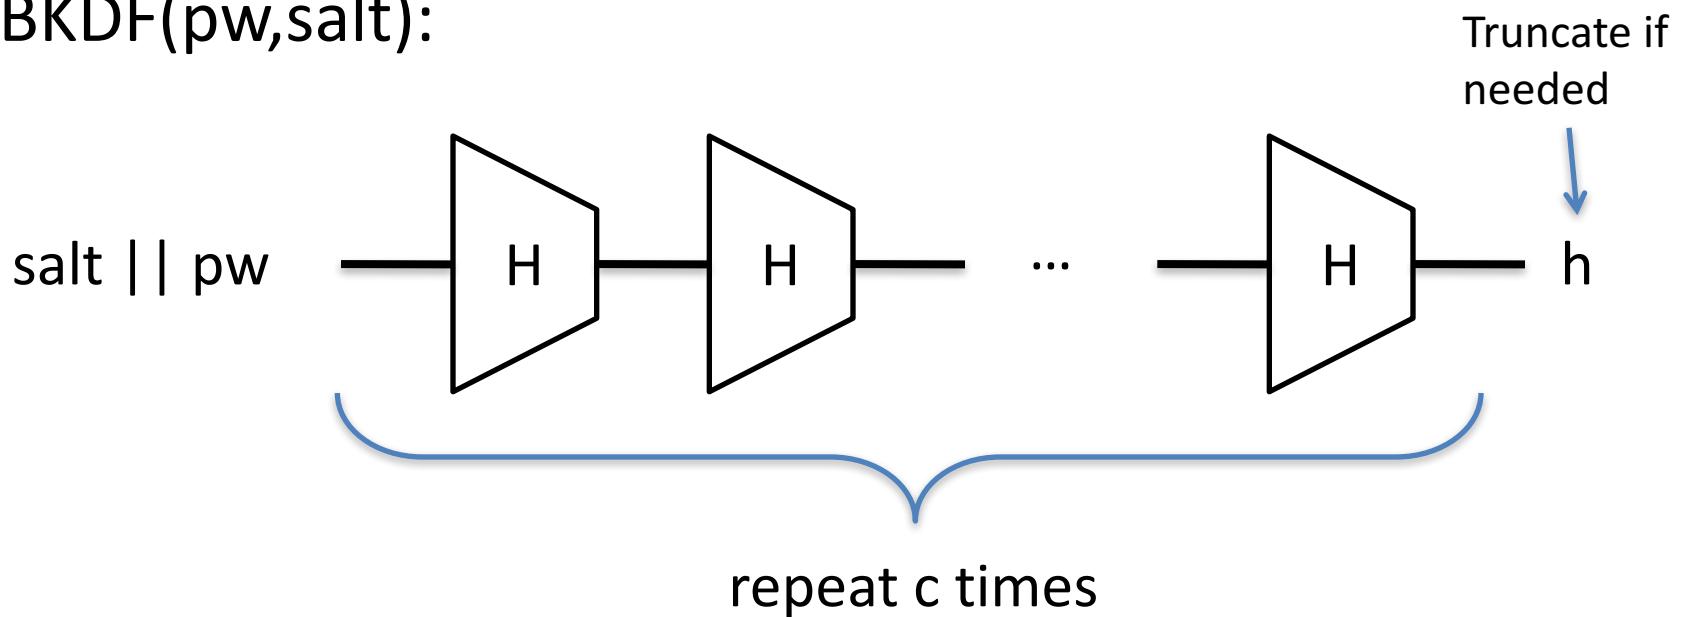
For each guess pw':

If  $H(\text{salt} \parallel \text{pw}') = h$  then  
Ret pw'

Rainbow tables speed this up in practice by way of precomputation. Large salts make rainbow tables impractical

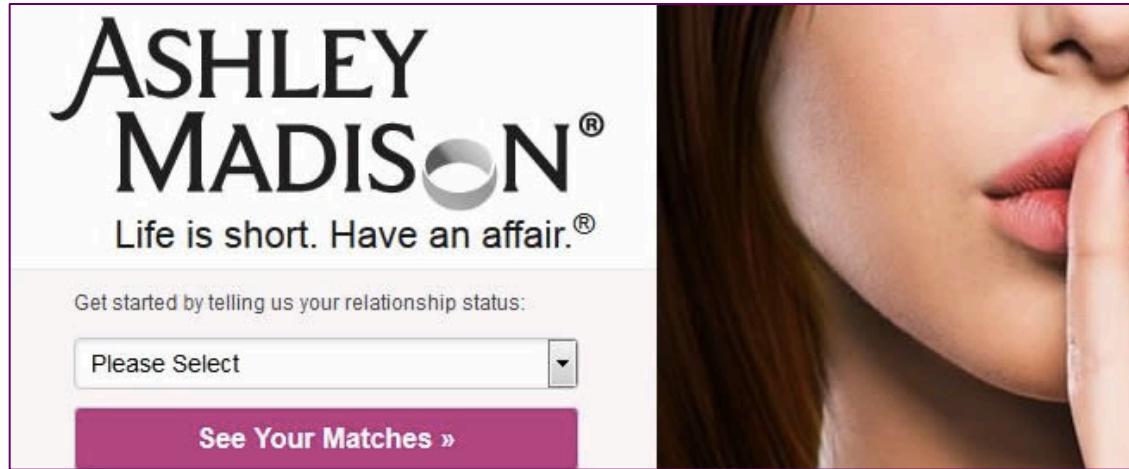
# Password-based Key Derivation (PBKDF)

PBKDF( $pw, salt$ ):



PKCS#5 standardizes PBKDF1 and PBKDF2, which are both hash-chain based. (PBKDF2 uses HMAC)

Slows down cracking attacks by a factor of  $c$



## AshleyMadison hack: 36 million user hashes

Salts + Passwords hashed using bcrypt with  $c = 2^{12} = 4096$   
4,007 cracked directly with trivial approach

CynoSure analysis: **11 million** hashes cracked  
>630,000 people used usernames as passwords  
MD5 hashes left lying around accidentally

```
ristenpart@Thomass-MacBook-Air:~/Dropbox/work/teaching/cs5830-spring2017/repo/slides$ openssl speed sha256
To get the most accurate results, try to run this
program when this computer is idle.
Doing sha256 for 3s on 16 size blocks: 4491209 sha256's in 2.98s
Doing sha256 for 3s on 64 size blocks: 2689214 sha256's in 2.98s
Doing sha256 for 3s on 256 size blocks: 1191470 sha256's in 2.99s
Doing sha256 for 3s on 1024 size blocks: 374944 sha256's in 2.98s
Doing sha256 for 3s on 8192 size blocks: 50404 sha256's in 2.99s
OpenSSL 0.9.8zg 14 July 2015
```

Say  $c = 4096$ . Generous back of envelope suggests that in 1 second, can test 367 passwords and so a naïve brute-force:

6 numerical digits	$10^6 =$ 1,000,000	~ 2724 seconds
6 lower case alphanumeric digits	$36^6 =$ 2,176,782,336	~ 68 days
8 alphanumeric + 10 special symbols	$72^8 =$ 722,204,136,308,736	~ 22 million days

Special-purpose hashing chips (built for Bitcoin)

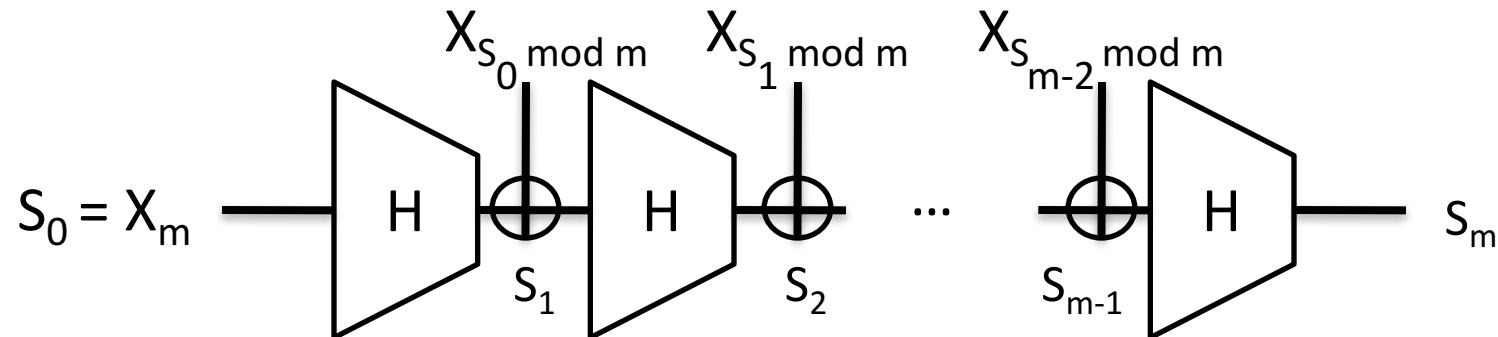
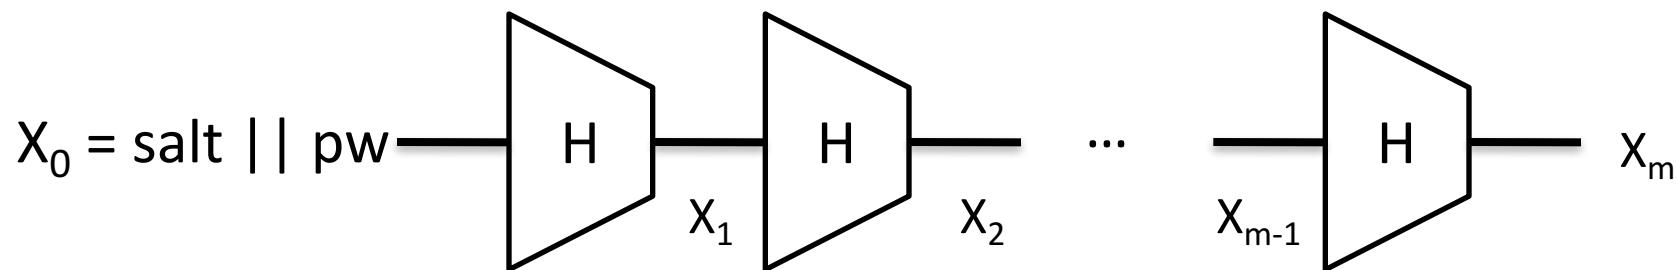
14,000,000,000 hashes / second

Can't be used easily for password cracking, but concern about ASICs (application-specific integrated circuits) remains



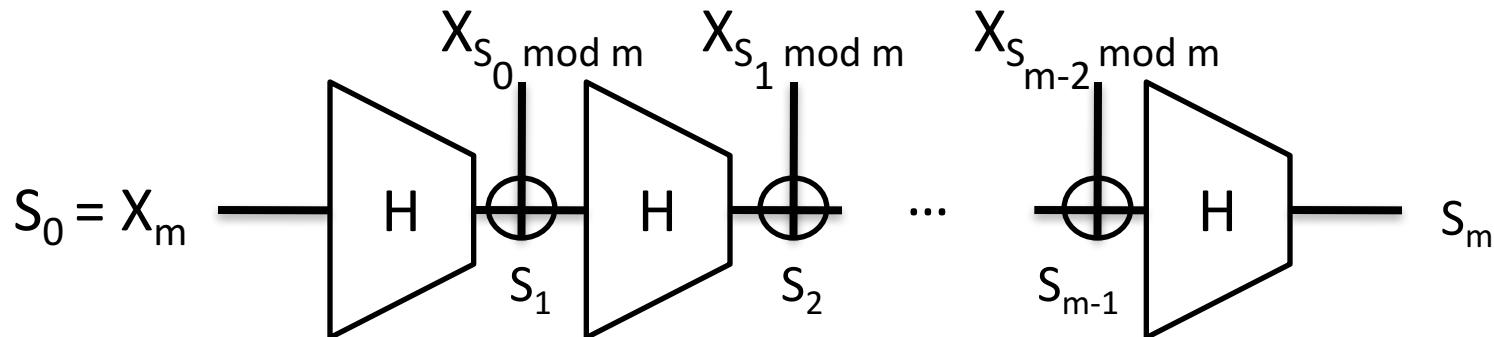
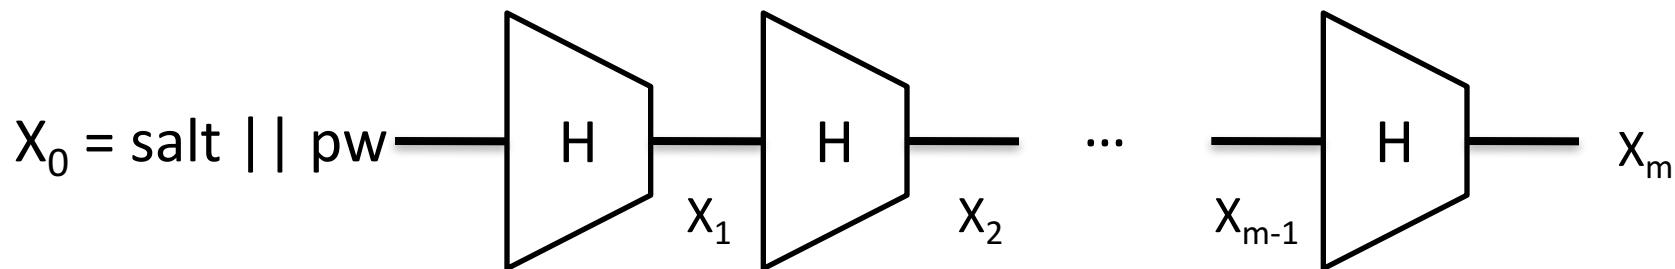
# Scrypt and memory-hard hashing

- Increase time & memory needed to compute hash.
  - This makes ASIC implementations harder
- Scrypt



# Scrypt and memory-hard hashing

- How much time-space required to compute?
  - Obvious algorithm:  $m$  space and  $2m$  time
  - Time-space complexity:  $O(m^2)$
  - Recent proof that no shortcut attacks exist [Alwen et al. 2016]



# Facebook password onion

```
$cur = 'password'  
$cur = md5($cur)  
$salt = randbytes(20)  
$cur = hmac_sha1($cur, $salt)  
$cur = remote_hmac_sha256($cur, $secret)  
$cur = scrypt($cur, $salt)  
$cur = hmac_sha256($cur, $salt)
```



# Strengthening password hash storage



tom, password1



h

f = HMAC(K, h)

$$h = H^c(\text{password1} \parallel \text{salt})$$

Store salt, f



Back-end  
crypto  
service



f = f'?

HMAC is pseudorandom function (PRF).

$$f' = H^c(123456 \parallel \text{salt})$$

f' = HMAC(K, h')

$$H^c(1234567 \parallel \text{salt})$$

$$H^c(12345 \parallel \text{salt})$$



Back-end  
crypto  
service

Must still perform online  
brute-force attack

Exfiltration doesn't help

# Facebook password onion

```
$cur = 'password'  
$cur = md5($cur)  
$salt = randbytes(20)  
$cur = hmac_sha1($cur, $salt)  
$cur = remote_hmac_sha256($cur, $secret)  
$cur = scrypt($cur, $salt)  
$cur = hmac_sha256($cur, $salt)
```



Evolution of their password hashing over time

*Limitations:*

- Can't rotate secret
- Can't do cryptographic erasure for compromise clean-up

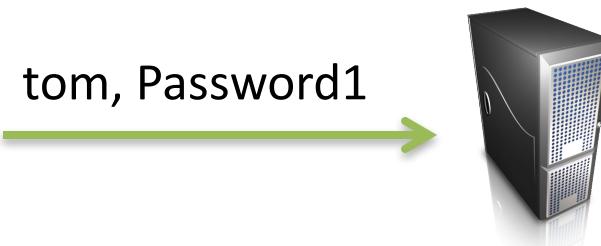
Pythia: better approach allowing secret key rotations

<http://pages.cs.wisc.edu/~ace/pythia.html>

# Simple typo-tolerant password checking



tom, Password1



tom	$G_K(\text{password1})$
alice	$G_K(123456)$
bob	$G_K(p@ssword!)$

Easily-corrected typos admit simple correctors:

$f_{\text{caps}}(x)$  =  $x$  with case of all letters flipped

$f_{\text{1st-case}}(x)$  =  $x$  with first letter case flipped, if it is letter

...

Define set C of corrector functions. Let  $h = G_K(\text{password1})$

Relaxed checker:

If  $G_K(\text{Password1}) = h$  then Return 1

For each  $f$  in C:

If  $G_K(f(\text{Password1})) = h$  then Return 1

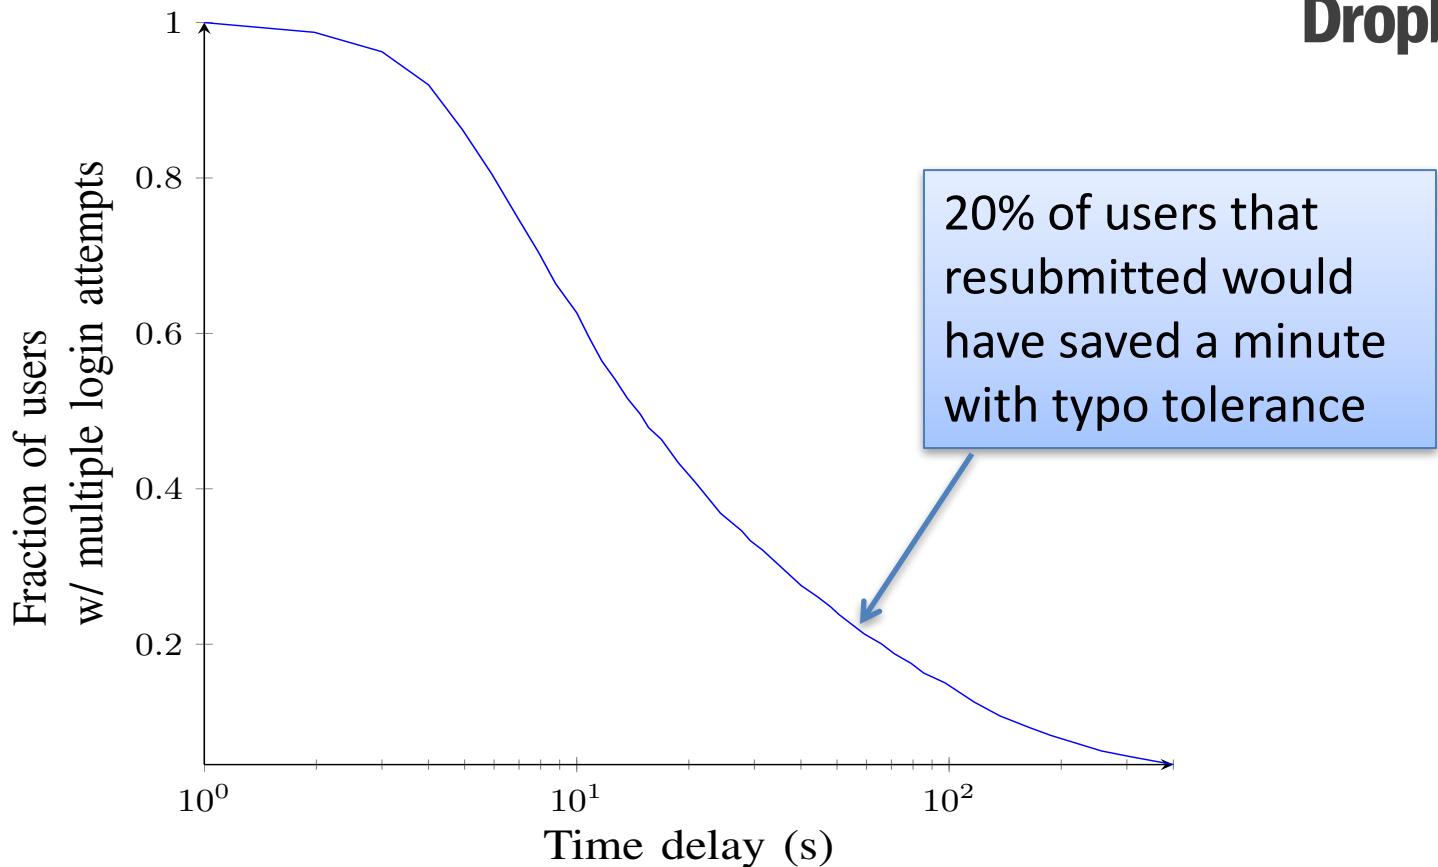
Return 0

$G_K$  is a password hashing scheme, possibly using secret K

# Dropbox experiments



Utility of caps lock, first letter capitalization, removing last character



Typo tolerance would add several person-months of login time

3% of *users* couldn't log in during 24 hour period, but could have with one of Top 3 correctors

# Empirical security analyses



Use password leaks as empirical password distributions

Typo-tolerant checker estimates distribution using RockYou

Simulate attacker following greedy strategy for typo-tolerant checker

Assume attacker knows challenge distribution exactly

Challenge Distribution	Correctors	Exact success	Greedy strategy
q = 100	phpbb	C <sub>top2</sub>	5.50%
	phpbb	C <sub>top3</sub>	5.50%
	Myspace	C <sub>top2</sub>	2.86%
	Myspace	C <sub>top3</sub>	2.86%

Attackers that estimate challenge distribution incorrectly:

Often perform worse when trying to take advantage of tolerance  
(See paper for details)

# TypTop: prototype adaptive checker

- Mechanical turk studies showed adaptivity can be beneficial
  - 45% of users would benefit from personalization
- We built a prototype called TypTop for OS X and Linux login
  - <https://typtop.info>

# Another application of PBKDFs: PW-based encryption

PWEnc(pw,M):

salt  $\leftarrow \$_{\{0,1\}^{256}}$

K  $\leftarrow \text{PBKDF}(\text{pw}, \text{salt})$

C  $\leftarrow \text{Enc}(K, M)$

Return (salt,C)

PWDec(pw,salt | | C):

K  $\leftarrow \text{PBKDF}(\text{pw}, \text{salt})$

M  $\leftarrow \text{Dec}(K, C)$

Return M

Enc is a *one-time-secure*  
AE scheme:

CTR-then-HMAC,  
constant IV for CTR

CBC-then-HMAC,  
constant IV for CBC mode

# What's wrong with this?

```
import base64
from cryptography.fernet import Fernet
from cryptography.hazmat.primitives import hashes
from cryptography.hazmat.backends import default_backend

def get_key(password):
    digest = hashes.Hash(hashes.SHA256(), backend=default_backend())
    digest.update(password)
    return base64.urlsafe_b64encode(digest.finalize())

def encrypt(password, token):
    f = Fernet(get_key(key))
    return f.encrypt(bytes(token))

def decrypt(password, token):
    f = Fernet(get_key(password))
    return f.decrypt(bytes(token))
```

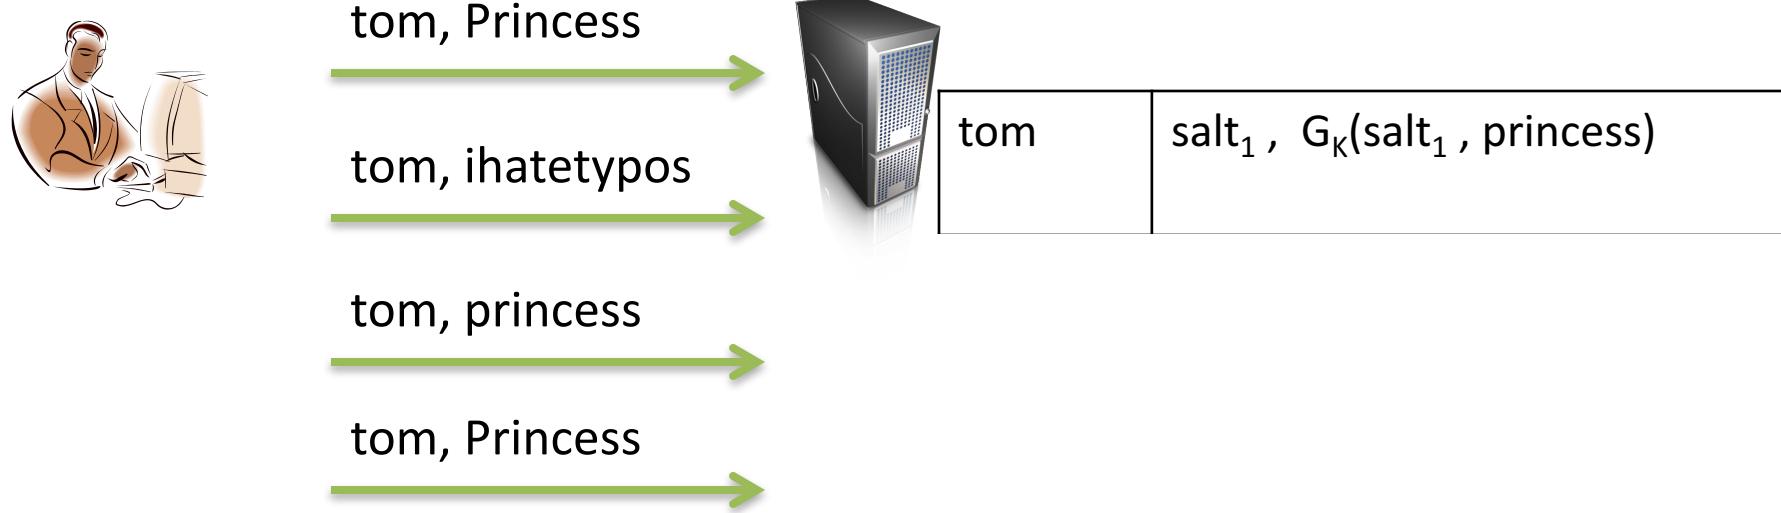
<http://incolumitas.com/2014/10/19/using-the-python-cryptography-module-with-custom-passwords/>

# Crypto library API challenges

- Tension between *opinionated API, security and application needs*
  - Opinionated:
    - Low-flexibility for caller (e.g., Fernet)
  - Un-opinionated:
    - High-flexibility for caller (e.g., Hazmat, OpenSSL)
- Opinionated best for security, but only if it handles application needs
  - Developers will work around your API, most likely insecurely!

# Personalized typo-tolerant checking

Another approach: learn typos individual user makes over time



Check  $G_K(salt_1, \text{Password1})$ , see that it is wrong

Add to a wait list of recent incorrect submissions

When user correctly logs in:

- Check wait list, apply valid typo policy (e.g., within edit distance 1 of true password)
- Add valid typos from wait list into cache
- Clear wait list

Check  $G_K(salt_1, \text{Password1})$  and  $G_K(salt_2, \text{Password1})$ , allow login if either match

# Personalized typo-tolerant checking

Another approach: learn typos individual user makes over time



tom, Princess  
tom, ihatetypos  
tom, princess  
tom, Princess

→  
→  
→  
→



tom	salt <sub>1</sub> , G <sub>K</sub> (salt <sub>1</sub> , princess)
Typo cache:	salt <sub>2</sub> , G <sub>K</sub> (salt <sub>2</sub> , Password1)
Wait list:	Password1    ihatetypos

Can't store wait list in clear, security problem!

# Personalized typo-tolerant checking

Another approach: learn typos individual user makes over time



tom, Princess  
tom, ihatetypos  
tom, princess  
tom, Princess

Four lines of text representing different typed versions of the password "tom". Each line has a green arrow pointing to the right, indicating they are being sent to a server.



$pk, E_{password1}(sk), E_{Password1}(sk)$	
tom	salt <sub>1</sub> , G <sub>K</sub> (salt <sub>1</sub> , princess)
Typo cache:	salt <sub>2</sub> , G <sub>K</sub> (salt <sub>2</sub> , Password1)
Wait list:	$Enc_{password1}(password1)$ $Enc_{pk}(ihatetypos)$

Obviously can't store wait list in clear, security problem

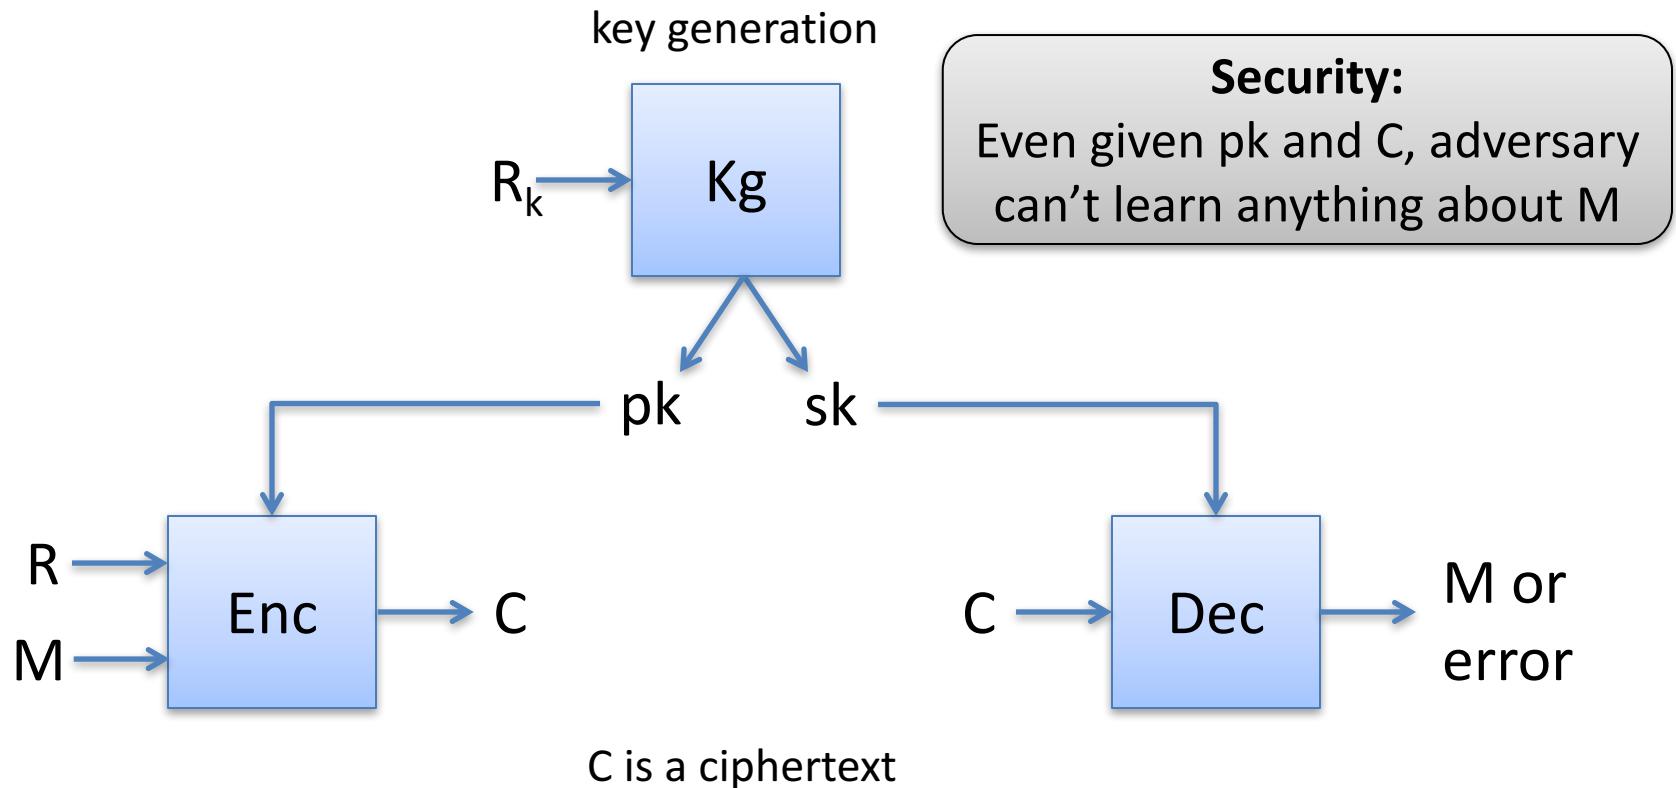
Can encrypt wait list using **public key encryption**

- Encrypt secret key at registration time using password1
- Encrypt secret key under each typo added to typo cache

Lots more details of design:

Randomizing order of typo cache, cache eviction policies, etc.

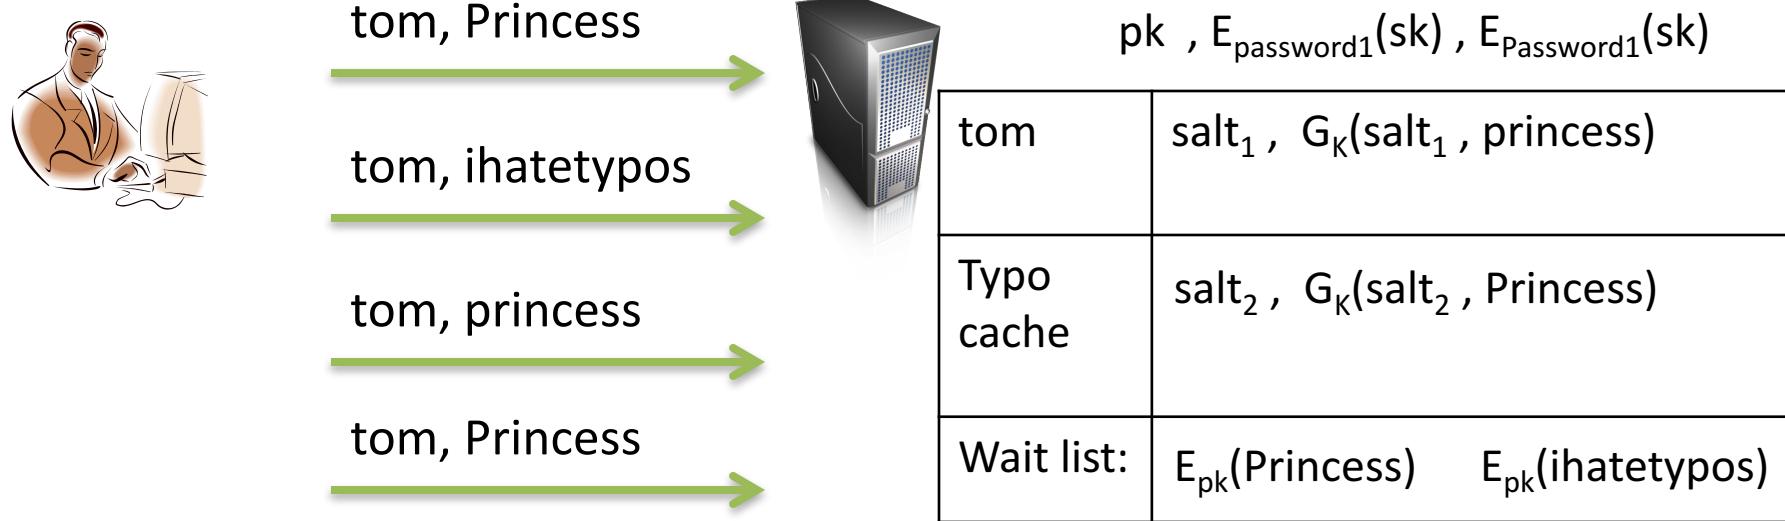
# Public-key encryption



Correctness:  $\text{Dec}(\text{sk}, \text{Enc}(\text{pk}, \text{M}, \text{R})) = \text{M}$  with probability 1 over randomness used

# Personalized typo-tolerant checking

Another approach: learn typos individual user makes over time



**Security:** we prove that for realistic password/typo distributions, an attacker that compromises system cannot do better than classic brute-force attack against  $G_K(\text{salt}_1 , \text{password1})$

No security loss by adding typo cache

# Summary

- Password hashing
  - Make hashing resource-intensive (time and/or memory)
  - Slows down brute-force cracking attacks
- Password-based authentication
  - Password onions
  - Typo-tolerance
- Password-based encryption